
SLAYER PyTorch

Release 0.1

Sumit Bam Shrestha

Aug 30, 2019

CONTENTS:

1	SLAYER PyTorch main	1
2	SLAYER module	3
3	SLAYER Loihi module	7
4	SLAYER Parameter	9
5	Spike Classifier	11
6	Spike Loss	13
7	Spike Input/Output	15
8	Learning statistics	21
9	Optimizer	25
10	Quantize module	27
11	Indices and tables	29
11.1	Usage:	29
11.2	Example:	29
	Python Module Index	31
	Index	33

SLAYER PYTORCH MAIN

This module bundles various SLAYER PyTorch modules as a single package. The complete module can be imported as

```
>>> import slayerSNN as snn
```

- The **slayer:spike-layer** module is available as `snn.layer`.
- The **slayer-Loihi:spike-layer** module is available as `snn.loihi`.
- The **yaml-parameter** module is available as `snn.params`.
- The **spike-loss** module is available as `snn.loss`.
- The **spike-classifier** module is available as `snn.predict`.
- The **spike-IO** module is available as `snn.io`.
- The **quantize** module is available as `snn.quantize`.

SLAYER MODULE

class `slayer.spikeLayer` (*neuronDesc*, *simulationDesc*, *fullRefKernel=False*)

This class defines the main engine of SLAYER. It provides necessary functions for describing a SNN layer. The input to output connection can be fully-connected, convolutional, or aggregation (pool) It also defines the psp operation and spiking mechanism of a spiking neuron in the layer.

Important: It assumes all the tensors that are being processed are 5 dimensional. (Batch, Channels, Height, Width, Time) or NCHWT format. The user must make sure that an input of correct dimension is supplied.

*If the layer does not have spatial dimension, the neurons can be distributed along either Channel, Height or Width dimension where Channel * Height * Width is equal to number of neurons. It is recommended (for speed reasons) to define the neuons in Channels dimension and make Height and Width dimension one.*

Arguments:

- **neuronDesc** (`slayerParams.yamlParams`): spiking neuron descriptor.

```
neuron:
  type:      SRMALPHA  # neuron type
  theta:     10        # neuron threshold
  tauSr:     10.0      # neuron time constant
  tauRef:    1.0       # neuron refractory time constant
  scaleRef:  2         # neuron refractory response scaling (relative to
↪theta)
  tauRho:    1         # spike function derivative time constant
↪(relative to theta)
  scaleRho:  1         # spike function derivative scale factor
```

- **simulationDesc** (`slayerParams.yamlParams`): simulation descriptor

```
simulation:
  Ts: 1.0          # sampling time (ms)
  tSample: 300     # time length of sample (ms)
```

- **fullRefKernel** (bool, optional): high resolution refractory kernel (the user shall not use it in practice)

Usage:

```
>>> snnLayer = slayer.spikeLayer(neuronDesc, simulationDesc)
```

conv (*inChannels*, *outChannels*, *kernelSize*, *stride=1*, *padding=0*, *dilation=1*, *groups=1*, *weightScale=100*)

Returns a function that can be called to apply conv layer mapping to input tensor per time instance. It behaves same as `torch.nn.conv2d` applied for each time instance.

Arguments:

- `inChannels (int)`: number of channels in input
- `outChannels (int)`: number of channels produced by convolution
- `kernelSize (int or tuple of two ints)`: size of the convolving kernel
- `stride (int or tuple of two ints)`: stride of the convolution. Default: 1
- `padding (int or tuple of two ints)`: zero-padding added to both sides of the input. Default: 0
- `dilation (int or tuple of two ints)`: spacing between kernel elements. Default: 1
- `groups (int or tuple of two ints)`: number of blocked connections from input channels to output channels. Default: 1
- `weightScale`: scale factor of default initialized weights. Default: 100

The parameters `kernelSize`, `stride`, `padding`, `dilation` can either be:

- a single `int` – in which case the same value is used for the height and width dimension
- a `tuple` of two `ints` – in which case, the first `int` is used for the height dimension, and the second `int` for the width dimension

Usage:

```
>>> conv = snnLayer.conv(2, 32, 5) # 32C5 filter
>>> output = conv(input)           # must have 2 channels
```

delay (*inputSize*)

Returns a function that can be called to apply delay operation in time dimension of the input tensor. The `delay` parameter is available as `delay.delay` and is initialized uniformly between 0ms and 1ms. The `delay` parameter is stored as float values, however, it is floored during actual delay application internally. The delay values are not clamped to zero. To maintain the causality of the network, one should clamp the delay values explicitly to ensure positive delays.

Arguments:

- `inputSize (int or tuple of three ints)`: spatial shape of the input signal in CHW format (Channel, Height, Width). If integer value is supplied, it refers to the number of neurons in channel dimension. Height and Width are assumed to be 1.

Usage:

```
>>> delay = snnLayer.delay((C, H, W))
>>> delayedSignal = delay(input)
```

Always clamp the delay after `optimizer.step()`.

```
>>> optimizer.step()
>>> delay.delay.data.clamp_(0)
```

delayShift (*input, delay, Ts=1*)

Applies delay in time dimension (assumed to be the last dimension of the tensor) of the input tensor. The autograd backward link is established as well.

Arguments:

- `input`: input Torch tensor.
- `delay (float or Torch tensor)`: amount of delay to apply. Same delay is applied to all the inputs if `delay` is `float` or Torch tensor of size 1. If the Torch tensor has size more than 1, its dimension must match the dimension of input tensor except the last dimension.

- Ts: sampling time of the delay. Default is 1.

Usage:

```
>>> delayedInput = slayer.delayShift(input, 5)
```

dense (*inFeatures*, *outFeatures*, *weightScale=10*)

Returns a function that can be called to apply dense layer mapping to input tensor per time instance. It behaves similar to `torch.nn.Linear` applied for each time instance.

Arguments:

- *inFeatures* (*int*, tuple of two ints, tuple of three ints): dimension of input features (Width, Height, Channel) that represents the number of input neurons.
- *outFeatures* (*int*): number of output neurons.
- *weightScale*: scale factor of default initialized weights. Default: 10

Usage:

```
>>> fcl = snnLayer.dense(2048, 512)           # takes (N, 2048, 1, 1, T) tensor
>>> fcl = snnLayer.dense((128, 128, 2), 512) # takes (N, 2, 128, 128, T)
↳ tensor
>>> output = fcl(input)                       # output will be (N, 512, 1, 1,
↳ T) tensor
```

dropout (*p=0.5*, *inplace=False*)

Returns a function that can be called to apply dropout layer to the input tensor. It behaves similar to `torch.nn.Dropout`. However, dropout over time dimension is preserved, i.e. if a neuron is dropped, it remains dropped for entire time duration.

Arguments:

- *p*: dropout probability.
- *inplace* (*bool*): inplace operation flag.

Usage:

```
>>> drop = snnLayer.dropout(0.2)
>>> output = drop(input)
```

pool (*kernelSize*, *stride=None*, *padding=0*, *dilation=1*)

Returns a function that can be called to apply pool layer mapping to input tensor per time instance. It behaves same as `torch.nn.sum` pooling applied for each time instance.

Arguments:

- *kernelSize* (*int* or tuple of two ints): the size of the window to pool over
- *stride* (*int* or tuple of two ints): stride of the window. Default: *kernelSize*
- *padding* (*int* or tuple of two ints): implicit zero padding to be added on both sides. Default: 0
- *dilation* (*int* or tuple of two ints): a parameter that controls the stride of elements in the window. Default: 1

The parameters *kernelSize*, *stride*, *padding*, *dilation* can either be:

- a single *int* – in which case the same value is used for the height and width dimension
- a tuple of two ints – in which case, the first *int* is used for the height dimension, and the second *int* for the width dimension

Usage:

```
>>> pool = snnLayer.pool(4) # 4x4 pooling
>>> output = pool(input)
```

psp (*spike*)

Applies psp filtering to spikes. The output tensor dimension is same as input.

Arguments:

- spike: input spike tensor.

Usage:

```
>>> filteredSpike = snnLayer.psp(spike)
```

pspFilter (*nFilter, filterLength, filterScale=1*)

Returns a function that can be called to apply a bank of temporal filters. The output tensor is of same dimension as input except the channel dimension is scaled by number of filters. The initial filters are initialized using default PyTorch initialization for conv layer. The filter banks are learnable. NOTE: the learned psp filter must be reversed because PyTorch performs correlation operation.

Arguments:

- nFilter: number of filters in the filterbank.
- filterLength: length of filter in number of time bins.
- filterScale: initial scaling factor for filter banks. Default: 1.

Usage:

```
>>> pspFilter = snnLayer.pspFilter()
>>> filteredSpike = pspFilter(spike)
```

pspLayer ()

Returns a function that can be called to apply psp filtering to spikes. The output tensor dimension is same as input. The initial psp filter corresponds to the neuron psp filter. The psp filter is learnable. NOTE: the learned psp filter must be reversed because PyTorch performs correlation operation.

Usage:

```
>>> pspLayer = snnLayer.pspLayer()
>>> filteredSpike = pspLayer(spike)
```

spike (*membranePotential*)

Applies spike function and refractory response. The output tensor dimension is same as input. *membranePotential* will reflect spike and refractory behaviour as well.

Arguments:

- membranePotential: subthreshold membrane potential.

Usage:

```
>>> outSpike = snnLayer.spike(membranePotential)
```

SLAYER LOIHI MODULE

class `slayerLoihi.spikeLayer` (*neuronDesc*, *simulationDesc*)

This class defines the main engine of SLAYER Loihi module. It is derived from `slayer.spikeLayer` with Loihi specific implementation for neuron model, weight quantization. All of the routines available for `slayer.spikeLayer` are applicable.

Arguments:

- **neuronDesc** (`slayerParams.yamlParams`): spiking neuron descriptor.

```
neuron:
  type:      LOIHI # neuron type
  vThMant:   80    # neuron threshold mantessa
  vDecay:    128   # compartment voltage decay
  iDecay:    1024  # compartment current decay
  refDelay:  1     # refractory delay
  wgtExp:    0     # weight exponent
  tauRho:    1     # spike function derivative time constant
  ↪(relative to theta)
  scaleRho:  1     # spike function derivative scale factor
```

- **simulationDesc** (`slayerParams.yamlParams`): simulation descriptor

```
simulation:
  Ts: 1.0          # sampling time (ms)
  tSample: 300     # time length of sample (ms)
```

Usage:

```
>>> snnLayer = slayerLoihi.spikeLayer(neuronDesc, simulationDesc)
```

conv (*inChannels*, *outChannels*, *kernelSize*, *stride=1*, *padding=0*, *dilation=1*, *groups=1*, *weightScale=100*, *quantize=True*)

This function behaves similar to `slayer.spikeLayer.conv()`. The only difference is that the weights are quantized with step of 2 (as is the case for signed weights in Loihi). One can, however, skip the quantization step altogether as well.

Arguments: The arguments that are different from `slayer.spikeLayer.conv()` are listed. *
weightScale: sale factor of default initialized weights. Default: 100 * *quantize* (bool): flag to quatize the weights or not. Default: True

Usage: Same as `slayer.spikeLayer.conv()`

dense (*inFeatures*, *outFeatures*, *weightScale=100*, *quantize=True*)

This function behaves similar to `slayer.spikeLayer.dense()`. The only difference is that the weights are quantized with step of 2 (as is the case for signed weights in Loihi). One can, however, skip the quantization step altogether as well.

Arguments: The arguments that are different from `slayer.spikeLayer.dense()` are listed. *
weightScale: scale factor of default initialized weights. Default: 100 * quantize (bool): flag
to quantize the weights or not. Default: True

Usage: Same as `slayer.spikeLayer.dense()`

pool (*kernelSize, stride=None, padding=0, dilation=1*)

This function behaves similar to `slayer.spikeLayer.pool()`. The only difference is that the weights are quantized with step of 2 (as is the case for signed weights in Loihi). One can, however, skip the quantization step altogether as well.

Arguments: The arguments set is same as `slayer.spikeLayer.pool()`.

Usage: Same as `slayer.spikeLayer.pool()`

spikeLoihi (*weightedSpikes*)

Applies Loihi neuron dynamics to weighted spike inputs and returns output spike tensor. The output tensor dimension is same as input.

NOTE: This function is different than the default `spike` function which takes membrane potential (weighted spikes with psp filter applied). Since the dynamics is modeled internally, it just takes in weightedSpikes (NOT FILTERED WITH PSP) for accurate Loihi neuron simulation.

Arguments:

- weightedSpikes: input spikes weighted by their corresponding synaptic weights.

Usage:

```
>>> outSpike = snnLayer.spikeLoihi(weightedSpikes)
```

spikeLoihiFull (*weightedSpikes*)

Applies Loihi neuron dynamics to weighted spike inputs and returns output spike, voltage and current. The output tensor dimension is same as input.

NOTE: This function does not have autograd routine in the computational graph.

Arguments:

- weightedSpikes: input spikes weighted by their corresponding synaptic weights.

Usage:

```
>>> outSpike, outVoltage, outCurrent = snnLayer.spikeLoihiFull(weightedSpikes)
```

SLAYER PARAMETER

This module provides a way to read the SLAYER configuration parameters from yaml file with dictionary like access. A typical yaml configuration file looks like this.

```
1 simulation:
2   Ts: 1.0
3   tSample: 1450
4 neuron:
5   type: SRMALPHA
6   theta: 10
7   tauSr: 1.0
8   tauRef: 1.0
9   scaleRef: 2      # relative to theta
10  tauRho: 1        # relative to theta
11  scaleRho: 1
12 layer:
13   - {dim: 34x34x2, wScale: 0.5}
14   - {dim: 16c5z}
15   - {dim: 2a}
16   - {dim: 64c3z}
17   - {dim: 2a}
18   - {dim: 512}
19   - {dim: 10}
20 training:
21   error:
22     type: NumSpikes #ProbSpikes #NumSpikes
23     tgtSpikeRegion: {start: 0, stop: 350}
24     tgtSpikeCount: {true: 60, false: 10}
25   path:
26     out: Trained/
27     in: path_to_spike_files
28     train: path_to_train_list
29     test: path_to_test_list
```

class slayerParams.yamlParams (parameter_file_path)

This class reads yaml parameter file and allows dictionary like access to the members.

Usage:

```
import slayerSNN as snn
netParams = snn.params('path_to_yaml_file')      # OR
netParams = yamlParams('path_to_yaml_file')

netParams['training']['learning']['etaW'] = 0.01
print('Simulation step size      ', netParams['simulation']['Ts'])
```

(continues on next page)

(continued from previous page)

```
print('Spiking neuron time constant', netParams['neuron']['tauSr'])
print('Spiking neuron threshold      ', netParams['neuron']['theta'])

netParams.save('filename.yaml')
```

SPIKE CLASSIFIER

class spikeClassifier.**spikeClassifier**

It provides classification modules for SNNs. All the functions it supplies are static and can be called without making an instance of the class.

static getClass (*spike*)

Returns the predicted class label. It assigns single class for the SNN output for the whole simulation runtime.

Usage:

```
>>> predictedClass = spikeClassifier.getClass(spikeOut)
```


SPIKE LOSS

class spikeLoss.**spikeLoss** (*networkDescriptor*, *slayerClass*=<class 'slayer.spikeLayer'>)

This class defines different spike based loss modules that can be used to optimize the SNN.

NOTE: By default, this class uses the spike kernels from `slayer.spikeLayer (snn.layer)`. In some cases, you may want to explicitly use different spike kernels, for e.g. `slayerLoihi.spikeLayer (snn.loihi)`. In that scenario, you can explicitly pass the class name: `slayerClass=snn.loihi`

Usage:

```
>>> error = spikeLoss.spikeLoss(networkDescriptor)
>>> error = spikeLoss.spikeLoss(errorDescriptor, neuronDesc, simulationDesc)
>>> error = spikeLoss.spikeLoss(netParams, slayerClass=slayerLoihi.spikeLayer)
```

numSpikes (*spikeOut*, *desiredClass*, *numSpikesScale*=1)

Calculates spike loss based on number of spikes within a *target region*. The *target region* and *desired spike count* is specified in `error.errorDescriptor['tgtSpikeRegion']`. Any spikes outside the target region are penalized with `error.spikeTime` loss..

$$e(t) = \begin{cases} \frac{actualSpikeCount - desiredSpikeCount}{targetRegionLength} & \text{for } t \in targetRegion \\ (\varepsilon * (output - desired))(t) & \text{otherwise} \end{cases}$$
$$E = \int_0^T e(t)^2 dt$$

Arguments:

- `spikeOut (torch.tensor)`: spike tensor
- `desiredClass (torch.tensor)`: one-hot encoded desired class tensor. Time dimension should be 1 and rest of the tensor dimensions should be same as `spikeOut`.

Usage:

```
>>> loss = error.numSpikes(spikeOut, target)
```

spikeTime (*spikeOut*, *spikeDesired*)

Calculates spike loss based on spike time. The loss is similar to van Rossum distance between output and desired spike train.

$$E = \int_0^T (\varepsilon * (output - desired))(t)^2 dt$$

Arguments:

- `spikeOut (torch.tensor)`: spike tensor
- `spikeDesired (torch.tensor)`: desired spike tensor

Usage:

```
>>> loss = error.spikeTime(spikeOut, spikeDes)
```

SPIKE INPUT/OUTPUT

`spikeFileIO.animTD (TD, fig=None, frameRate=24, preComputeFrames=True, repeat=True)`

Return animation object for TD event.

Arguments:

- `TD`: spike event to visualize.
- `fig`: figure to plot animation. Default is `None`, in which case a figure is created.
- `frameRate`: framerate of visualization.
- `preComputeFrames`: flag to enable precomputation of frames for faster visualization. Default is `True`.
- `repeat`: flag to enable repeat of animation. Default is `True`.

Usage:

```
>>> anim = animTD(TD)
```

`spikeFileIO.encode1DnumSpikes (filename, nID, tSt, tEn, nSp)`

Writes binary spike file given a tuple specifying neuron, start of spike region, end of spike region and number of spikes.

The binary file is encoded as follows:

- Number of spikes data is represented by an 80 bit number
- First 16 bits (bits 79-64) represent the neuronID
- Next 24 bits (bits 63-40) represents the start time in microseconds
- Next 24 bits (bits 39-16) represents the end time in microseconds
- Last 16 bits (bits 15-0) represents the number of spikes

Arguments:

- `filename (string)`: path to the binary file
- `nID (numpy array)`: neuron ID
- `tSt (numpy array)`: region start time (in milliseconds)
- `tEn (numpy array)`: region end time (in milliseconds)
- `nSp (numpy array)`: number of spikes in the region

Usage:

```
>>> spikeFileIO.encode1DnumSpikes(file_path, nID, tSt, tEn, nSp)
```

`spikeFileIO.encode1Dspikes` (*filename*, *TD*)

Writes one dimensional binary spike file from a TD event.

The binary file is encoded as follows:

- Each spike event is represented by a 40 bit number.
- First 16 bits (bits 39-24) represent the neuronID.
- Bit 23 represents the sign of spike event: 0=>OFF event, 1=>ON event.
- the last 23 bits (bits 22-0) represent the spike event timestamp in microseconds.

Arguments:

- `filename` (string): path to the binary file.
- `TD` (an `spikeFileIO.event`): TD event.

Usage:

```
>>> spikeFileIO.write1Dspikes(file_path, TD)
```

`spikeFileIO.encode2Dspikes` (*filename*, *TD*)

Writes two dimensional binary spike file from a TD event. It is the same format used in neuromorphic datasets NMNIST & NCALTECH101.

The binary file is encoded as follows:

- Each spike event is represented by a 40 bit number.
- First 8 bits (bits 39-32) represent the xID of the neuron.
- Next 8 bits (bits 31-24) represent the yID of the neuron.
- Bit 23 represents the sign of spike event: 0=>OFF event, 1=>ON event.
- The last 23 bits (bits 22-0) represent the spike event timestamp in microseconds.

Arguments:

- `filename` (string): path to the binary file.
- `TD` (an `spikeFileIO.event`): TD event.

Usage:

```
>>> spikeFileIO.write2Dspikes(file_path, TD)
```

`spikeFileIO.encode3Dspikes` (*filename*, *TD*)

Writes binary spike file for TD event in height, width and channel dimension.

The binary file is encoded as follows:

- Each spike event is represented by a 56 bit number.
- First 12 bits (bits 56-44) represent the xID of the neuron.
- Next 12 bits (bits 43-32) represent the yID of the neuron.
- Next 8 bits (bits 31-24) represents the channel ID of the neuron.
- The last 24 bits (bits 23-0) represent the spike event timestamp in microseconds.

Arguments:

- filename (string): path to the binary file.
- TD (an spikeFileIO.event): TD event.

Usage:

```
>>> spikeFileIO.write3Dspikes(file_path, TD)
```

spikeFileIO.encodeNpSpikes (filename, TD, fmt='xypt', timeUnit=0.001)
Writes TD event into numpy file.

Arguments:

- filename (string): path to the binary file.
- TD (an spikeFileIO.event): TD event.

Usage:

```
>>> spikeFileIO.write1Dspikes(file_path, TD)
>>> spikeFileIO.write1Dspikes(file_path, TD, fmt='xypt')
```

class spikeFileIO.event (xEvent, yEvent, pEvent, tEvent)
This class provides a way to store, read, write and visualize spike event.

Members:

- x (numpy int array): *x* index of spike event.
- y (numpy int array): *y* index of spike event (not used if the spatial dimension is 1).
- p (numpy int array): *polarity* or *channel* index of spike event.
- t (numpy double array): *timestamp* of spike event. Time is assumed to be in ms.

Usage:

```
>>> TD = spikeFileIO.event(xEvent, yEvent, pEvent, tEvent)
```

toSpikeArray (samplingTime=1, dim=None)

Returns a numpy tensor that contains the spike events sampled in bins of *samplingTime*. The array is of dimension (channels, height, time) or “CHT” for 1D data. The array is of dimension (channels, height, width, time) or “CHWT” for 2D data.

Arguments:

- samplingTime: the width of time bin to use.
- dim: the dimension of the desired tensor. Assigns dimension itself if not provided.

Usage:

```
>>> spike = TD.toSpikeArray()
```

toSpikeTensor (emptyTensor, samplingTime=1, randomShift=False)

Returns a numpy tensor that contains the spike events sampled in bins of *samplingTime*. The tensor is of dimension (channels, height, width, time) or “CHWT”.

Arguments:

- emptyTensor (numpy or torch tensor): an empty tensor to hold spike data
- samplingTime: the width of time bin to use.
- randomShift: flag to shift the sample in time or not. Default: False.

Usage:

```
>>> spike = TD.toSpikeTensor( torch.zeros((2, 240, 180, 5000)) )
```

`spikeFileIO.read1DnumSpikes(filename)`

Reads a tuple specifying neuron, start of spike region, end of spike region and number of spikes from binary spike file.

The binary file is encoded as follows:

- Number of spikes data is represented by an 80 bit number.
- First 16 bits (bits 79-64) represent the neuronID.
- Next 24 bits (bits 63-40) represents the start time in microseconds.
- Next 24 bits (bits 39-16) represents the end time in microseconds.
- Last 16 bits (bits 15-0) represents the number of spikes.

Arguments:

- `filename (string)`: path to the binary file

Usage:

```
>>> nID, tSt, tEn, nSp = spikeFileIO.read1DnumSpikes(file_path)
``tSt`` and ``tEn`` are returned in milliseconds
```

`spikeFileIO.read1Dspikes(filename)`

Reads one dimensional binary spike file and returns a TD event.

The binary file is encoded as follows:

- Each spike event is represented by a 40 bit number.
- First 16 bits (bits 39-24) represent the neuronID.
- Bit 23 represents the sign of spike event: 0=>OFF event, 1=>ON event.
- the last 23 bits (bits 22-0) represent the spike event timestamp in microseconds.

Arguments:

- `filename (string)`: path to the binary file.

Usage:

```
>>> TD = spikeFileIO.read1Dspikes(file_path)
```

`spikeFileIO.read2Dspikes(filename)`

Reads two dimensional binary spike file and returns a TD event. It is the same format used in neuromorphic datasets NMNIST & NCALTECH101.

The binary file is encoded as follows:

- Each spike event is represented by a 40 bit number.
- First 8 bits (bits 39-32) represent the xID of the neuron.
- Next 8 bits (bits 31-24) represent the yID of the neuron.
- Bit 23 represents the sign of spike event: 0=>OFF event, 1=>ON event.
- The last 23 bits (bits 22-0) represent the spike event timestamp in microseconds.

Arguments:

- `filename (string)`: path to the binary file.

Usage:

```
>>> TD = spikeFileIO.read2Dspikes(file_path)
```

`spikeFileIO.read3Dspikes (filename)`

Reads binary spike file for spike event in height, width and channel dimension and returns a TD event.

The binary file is encoded as follows:

- Each spike event is represented by a 56 bit number.
- First 12 bits (bits 56-44) represent the xID of the neuron.
- Next 12 bits (bits 43-32) represent the yID of the neuron.
- Next 8 bits (bits 31-24) represents the channel ID of the neuron.
- The last 24 bits (bits 23-0) represent the spike event timestamp in microseconds.

Arguments:

- `filename (string)`: path to the binary file.

Usage:

```
>>> TD = spikeFileIO.read3Dspikes(file_path)
```

`spikeFileIO.readNpSpikes (filename, fmt='xypt', timeUnit=0.001)`

Reads numpy spike event and returns a TD event. The numpy array is assumed to be of nEvent x event diension.

Arguments:

- `filename (string)`: path to the file.
- `fmt (string)`: format of event. For e.g. 'xypt' means the event data is arrange in x data, y data, p data and time data.
- `timeUnit (double)`: factor to scale the time data to convert it into seconds. Default: 1e-3 (ms).

Usage:

```
>>> TD = spikeFileIO.readNpSpikes(file_path)
>>> TD = spikeFileIO.readNpSpikes(file_path, fmt='xypt')
>>> TD = spikeFileIO.readNpSpikes(file_path, timeUnit=1e-6)
```

`spikeFileIO.showTD (TD, fig=None, frameRate=24, preComputeFrames=True, repeat=False)`

Visualizes TD event.

Arguments:

- `TD`: spike event to visualize.
- `fig`: figure to plot animation. Default is None, in which case a figure is created.
- `frameRate`: framerate of visualization.
- `preComputeFrames`: flag to enable precomputation of frames for faster visualization. Default is True.
- `repeat`: flag to enable repeat of animation. Default is False.

Usage:

```
>>> showTD (TD)
```

`spikeFileIO.spikeArrayToEvent (spikeMat, samplingTime=1)`

Returns TD event from a numpy array (of dimension 3 or 4). The numpy array must be of dimension (channels, height, time) or “CHT” for 1D data. The numpy array must be of dimension (channels, height, width, time) or “CHWT” for 2D data.

Arguments:

- `spikeMat`: numpy array with spike information.
- `samplingTime`: time width of each time bin.

Usage:

```
>>> TD = spikeFileIO.spikeArrayToEvent (spike)
```


LEARNING STATISTICS

class learningStats.learningStat

This class collect the learning statistics over the epoch.

Usage:

This class is designed to be used with learningStats instance although it can be used separately.

```
>>> trainingStat = learningStat()
```

accuracy ()

Returns the average accuracy calculated from the point the stats was reset.

Usage:

```
>>> accuracy = trainingStat.accuracy()
```

loss ()

Returns the average loss calculated from the point the stats was reset.

Usage:

```
>>> loss = trainingStat.loss()
```

reset ()

Reset the learning staistics. This should usually be done before the start of an epoch so that new statistics counts can be accumulated.

Usage:

```
>>> trainingStat.reset()
```

update ()

Updates the stats of the current session and resets the measures for next session.

Usage:

```
>>> trainingStat.update()
```

class learningStats.learningStats

This class provides mechanism to collect learning stats for training and testing, and displaying them efficiently.

Usage:

```
stats = learningStats()
```

(continues on next page)

(continued from previous page)

```

for epoch in range(100):
    tSt = datetime.now()

    stats.training.reset()
    for i in trainingLoop:
        # other main stuffs
        stats.training.correctSamples += numberOfCorrectClassification
        stats.training.numSamples += numberOfSamplesProcessed
        stats.training.lossSum += currentLoss
        stats.print(epoch, i, (datetime.now() - tSt).total_seconds())
    stats.training.update()

    stats.testing.reset()
    for i in testingLoop:
        # other main stuffs
        stats.testing.correctSamples += numberOfCorrectClassification
        stats.testing.numSamples += numberOfSamplesProcessed
        stats.testing.lossSum += currentLoss
        stats.print(epoch, i)
    stats.training.update()

```

load (filename="", numEpoch=None, modulo=1)

Loads the learning statistics logs from saved files.

Arguments:

- filename: filename to save the logs. accuracy.txt and loss.txt will be appended.
- numEpoch: number of epochs of logs to load. Default: None. numEpoch will be automatically determined from saved files.
- modulo: the gap in number of epoch before model was saved.

Usage:

```

# save stats
stats.load(epoch=10)

# save stats filename specified
stats.save(filename='Run101-0.001-', epoch=50) # Run101-0.001-accuracy.txt
↪and Run101-0.001-loss.txt

```

plot (figures=(1, 2), saveFig=False, path="")

Plots the available learning statistics.

Arguments:

- figures: Index of figure ID to plot on. Default is figure(1) for loss plot and figure(2) for accuracy plot.
- saveFig``(``bool): flag to save figure into a file.
- path: path to save the file. Defaule is ' '.

Usage:

```

# plot stats
stats.plot()

```

(continues on next page)

(continued from previous page)

```
# plot stats figures specified
stats.print(figures=(10, 11))
```

print (*epoch*, *iter=None*, *timeElapsed=None*)

Prints the available learning statistics from the current session on the console. For Linux systems, prints the data on same terminal space (might not work properly on other systems).

Arguments:

- *epoch*: epoch counter to display (required).
- *iter*: iteration counter to display (not required).
- *timeElapsed*: runtime information (not required).

Usage:

```
# prints stats with epoch index provided
stats.print(epoch)

# prints stats with epoch index and iteration index provided
stats.print(epoch, iter=i)

# prints stats with epoch index, iteration index and time elapsed information_
↪ provided
stats.print(epoch, iter=i, timeElapsed=time)
```

save (*filename=""*)

Saves the learning statistics logs.

Arguments:

- *filename*: filename to save the logs. `accuracy.txt` and `loss.txt` will be appended.

Usage:

```
# save stats
stats.save()

# save stats filename specified
stats.save(filename='Run101-0.001-') # Run101-0.001-accuracy.txt and Run101-0.
↪ 001-loss.txt
```

update ()

Updates the stats for training and testing and resets the measures for next session.

Usage:

```
>>> stats.update()
```


OPTIMIZER

class optimizer.**Nadam**(*params*, *lr*=0.001, *betas*=(0.9, 0.999), *eps*=1e-08, *weight_decay*=0, *amsgrad*=False)

Implements Nadam algorithm. (Modified Adam from [PyTorch](#))

It has been proposed in [Incorporating Nesterov Momentum into Adam](#).

Arguments:

- *params* (iterable): iterable of parameters to optimize or dicts defining parameter groups.
- *lr* (float, optional): learning rate (default: 1e-3).
- *betas* (Tuple[float, float], optional): coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999)).
- *eps* (float, optional): term added to the denominator to improve numerical stability (default: 1e-8).
- *weight_decay* (float, optional): weight decay (L2 penalty) (default: 0).
- *amsgrad* (boolean, optional): whether to use the AMSGrad variant of this algorithm from the paper [On the Convergence of Adam and Beyond](#) (default: False).

step (*closure*=None)

Performs a single optimization step.

Arguments:

- **closure** (callable, optional): A closure that reevaluates the model and returns the loss.

QUANTIZE MODULE

class quantizeParams.quantizeWeights

This class provides routine to quantize the weights during forward propagation pipeline. The backward propagation pipeline passes the gradient as it is, without any modification.

Arguments;

- **weights**: full precision weight tensor.
- **step**: quantization step size. Default: 1

Usage:

```
>>> # Quantize weights in step of 0.5
>>> stepWeights = quantizeWeights.apply(fullWeights, step=0.5)
```

static backward (*ctx, gradOutput*)

static forward (*ctx, weights, step=1*)

INDICES AND TABLES

- `genindex`
- `modindex`
- `search`

11.1 Usage:

```
>>> import slayerSNN as snn
```

- The **slayer:spike-layer** module is available as `snn.layer`.
- The **slayer-Loihi:spike-layer** module is available as `snn.loihi`
- The **yaml-parameter** module is available as `snn.params`.
- The **spike-loss** module is available as `snn.loss`.
- The **spike-classifier** module is available as `snn.predict`.
- The **spike-IO** module is available as `snn.io`.
- The **quantize** module is available as `snn.quantize`.

11.2 Example:

The SNN parameters are stored in a yaml file. The structure of the yaml file follows the same hierarchy as the **C++ SLAYER framework** (see Network Description)

A set of working example for SLAYER are available at [SLAYER PyTorch repo: examples](#)

A set of working example for SLAYER Loihi are available at [SLAYER PyTorch repo: examplesLoihi](#)

```
1 import slayerSNN as snn
2 # other imports and definitions
3
4 class Network(torch.nn.Module):
5     def __init__(self, netParams, device=device):
6         super(Network, self).__init__()
7         # initialize slayer
8         slayer = snn.layer(netParams['neuron'], netParams['simulation'], device=device)
9         self.sl = slayer
10        # define network functions
```

(continues on next page)

(continued from previous page)

```

11     self.conv1 = slayer.conv(2, 16, 5, padding=1)
12     self.conv2 = slayer.conv(16, 32, 3, padding=1)
13     self.conv3 = slayer.conv(32, 64, 3, padding=1)
14     self.pool1 = slayer.pool(2)
15     self.pool2 = slayer.pool(2)
16     self.fcl   = slayer.dense((8, 8, 64), 10)
17
18     def forward(self, spikeInput):
19         spikeLayer1 = self.sl.spike(self.conv1(self.sl.psp(spikeInput))) # 32, 32, 16
20         spikeLayer2 = self.sl.spike(self.pool1(self.sl.psp(spikeLayer1))) # 16, 16, 16
21         spikeLayer3 = self.sl.spike(self.conv2(self.sl.psp(spikeLayer2))) # 16, 16, 32
22         spikeLayer4 = self.sl.spike(self.pool2(self.sl.psp(spikeLayer3))) # 8, 8, 32
23         spikeLayer5 = self.sl.spike(self.conv3(self.sl.psp(spikeLayer4))) # 8, 8, 64
24         spikeOut    = self.sl.spike(self.fcl   (self.sl.psp(spikeLayer5))) # 10
25         return spikeOut
26
27     # network
28     net = Network(snn.params('path to yaml file'))
29
30     # cost function
31     error = snn.loss(netParams)
32
33     # dataloader not shown. input and target are assumed to be available
34     output = net.forward(input)
35     loss = error.numSpikes(output, target)

```

Important: It is assumed that all the tensors that are being processed are 5 dimensional. (Batch, Channels, Height, Width, Time) or NCHWT format. The user must make sure that an input of correct dimension is supplied.

*If the layer does not have spatial dimension, the neurons can be distributed along either Channel, Height or Width dimension where Channel * Height * Width is equal to number of neurons. It is recommended (for speed reasons) to define the neurons in Channels dimension and make Height and Width dimension one.*

PYTHON MODULE INDEX

I

`learningStats`, [21](#)

O

`optimizer`, [25](#)

Q

`quantizeParams`, [27](#)

S

`slayer`, [3](#)

`slayerLoihi`, [7](#)

`slayerParams`, [9](#)

`slayerSNN`, [1](#)

`spikeClassifier`, [11](#)

`spikeFileIO`, [15](#)

`spikeLoss`, [13](#)

A

`accuracy()` (*learningStats.learningStat method*), 21
`animTD()` (*in module spikeFileIO*), 15

B

`backward()` (*quantizeParams.quantizeWeights static method*), 27

C

`conv()` (*slayer.spikeLayer method*), 3
`conv()` (*slayerLoihi.spikeLayer method*), 7

D

`delay()` (*slayer.spikeLayer method*), 4
`delayShift()` (*slayer.spikeLayer method*), 4
`dense()` (*slayer.spikeLayer method*), 5
`dense()` (*slayerLoihi.spikeLayer method*), 7
`dropout()` (*slayer.spikeLayer method*), 5

E

`encode1DnumSpikes()` (*in module spikeFileIO*), 15
`encode1Dspikes()` (*in module spikeFileIO*), 16
`encode2Dspikes()` (*in module spikeFileIO*), 16
`encode3Dspikes()` (*in module spikeFileIO*), 16
`encodeNpSpikes()` (*in module spikeFileIO*), 17
`event` (*class in spikeFileIO*), 17

F

`forward()` (*quantizeParams.quantizeWeights static method*), 27

G

`getClass()` (*spikeClassifier.spikeClassifier static method*), 11

L

`learningStat` (*class in learningStats*), 21
`learningStats` (*class in learningStats*), 21
`learningStats` (*module*), 21
`load()` (*learningStats.learningStats method*), 22
`loss()` (*learningStats.learningStat method*), 21

N

`Nadam` (*class in optimizer*), 25
`numSpikes()` (*spikeLoss.spikeLoss method*), 13

O

`optimizer` (*module*), 25

P

`plot()` (*learningStats.learningStats method*), 22
`pool()` (*slayer.spikeLayer method*), 5
`pool()` (*slayerLoihi.spikeLayer method*), 8
`print()` (*learningStats.learningStats method*), 23
`psp()` (*slayer.spikeLayer method*), 6
`pspFilter()` (*slayer.spikeLayer method*), 6
`pspLayer()` (*slayer.spikeLayer method*), 6

Q

`quantizeParams` (*module*), 27
`quantizeWeights` (*class in quantizeParams*), 27

R

`read1DnumSpikes()` (*in module spikeFileIO*), 18
`read1Dspikes()` (*in module spikeFileIO*), 18
`read2Dspikes()` (*in module spikeFileIO*), 18
`read3Dspikes()` (*in module spikeFileIO*), 19
`readNpSpikes()` (*in module spikeFileIO*), 19
`reset()` (*learningStats.learningStat method*), 21

S

`save()` (*learningStats.learningStats method*), 23
`showTD()` (*in module spikeFileIO*), 19
`slayer` (*module*), 3
`slayerLoihi` (*module*), 7
`slayerParams` (*module*), 9
`slayerSNN` (*module*), 1
`spike()` (*slayer.spikeLayer method*), 6
`spikeArrayToEvent()` (*in module spikeFileIO*), 20
`spikeClassifier` (*class in spikeClassifier*), 11
`spikeClassifier` (*module*), 11
`spikeFileIO` (*module*), 15
`spikeLayer` (*class in slayer*), 3

`spikeLayer` (*class in `slayerLoihi`*), 7
`spikeLoihi()` (*`slayerLoihi.spikeLayer` method*), 8
`spikeLoihiFull()` (*`slayerLoihi.spikeLayer`
method*), 8
`spikeLoss` (*class in `spikeLoss`*), 13
`spikeLoss` (*module*), 13
`spikeTime()` (*`spikeLoss.spikeLoss` method*), 13
`step()` (*`optimizer.Nadam` method*), 25

T

`toSpikeArray()` (*`spikeFileIO.event` method*), 17
`toSpikeTensor()` (*`spikeFileIO.event` method*), 17

U

`update()` (*`learningStats.learningStat` method*), 21
`update()` (*`learningStats.learningStats` method*), 23

Y

`yamlParams` (*class in `slayerParams`*), 9